

# Pareto-Optimal Algorithms for Learning in Repeated Games

Talk at Central Applied Science, Meta



Natalie Collina, Upenn



Jon Schneider, Google Research




# Motivation

Agents use Learning Algorithms to make Decisions





# Motivation

Agents use Learning Algorithms to play games




# Motivation

Agents use Learning Algorithms to play repeated games



.....



# Motivation

## Repeated Ad-Auctions

We use an ad auction to determine the best ad to show to a person at a given point in time. The winning ad maximizes value for both people and businesses. Understanding the ad auction can help you understand your ad performance.

### When do ad auctions take place?

Each time there's an opportunity to show an ad to someone, an auction takes place to determine which ad to show to that person. Billions of auctions take place everyday across the Facebook family of apps.




### Who competes in each auction?

When advertisers create ads, they tell us who they want to show their ads to by defining a target audience. A person can fall into multiple target audiences. For example, one advertiser targets women who like skiing, while another advertiser targets all skiers who live in California. The same person (in this case, a female skier who lives in California) could fall into the target audience of both advertisers.


When there's an opportunity to show someone an ad, the ads with a target audience that the person belongs to are eligible to compete in the auction.

# Motivation

## Repeated Ad-Auctions



Automated  
Auctioneer





# Some Questions

- What are good learning algorithms to use?
- Existing Benchmark : No-Regret
- A New Criterion: Non-Manipulability
- Our Novel Criterion : Pareto-Optimality
  
- How might other agents respond to these learning algorithms?
- For eg: How should an auctioneer pick a dynamic pricing rule against certain bidding algorithms?

# Model

# Model

## Two Players - Learner and Optimizer




The Learner knows their own payoff but not that of the optimizer




The optimizer has full information and best-responds (non-myopically)

# Model


## Two Players - Learner and Optimizer



The Learner knows their own payoff but not that of the optimizer



The learner observes the action played by the optimizer in each round



Two-player bimatrix game

(Repeated T times)





The optimizer has full information and best-responds (non-myopically)

# Example - The RPS game

A Two-Player Zero-sum Game




			
	0	-1	1
	1	0	-1
	-1	1	0

Learner Payoffs

Unknown to the Learner



		
0	1	-1
-1	0	1
1	-1	0

Optimizer Payoffs

# Example - The RPS game



0	-1	1
1	0	-1
-1	1	0

Learner Payoffs



Unknown to the Learner



1	0	0
0	1	0
0	0	1

Optimizer Payoffs

# Model: Notation




## In Each Round


- The learner has action set  $\Delta_m$
- The optimizer has action set  $\Delta_n$
- They play actions  $y_t, x_t$  in the t-th round.
- Linear utility function  $u_L, u_O$

# Model : Learning Algorithms

## The Learner Perspective



Without knowing  $u_O$ , the learner commit to an algorithm mapping (deterministically) from histories of play of length  $t-1$  to a distribution  $y_t$  over actions in the  $t$ -th round



# Model

## The Optimizer Perspective



With full information (payoffs, learner algorithm), the optimizer plays a best-response sequence of actions

$$x_1, x_2 \dots x_T \in \operatorname{argmax}_{(x_1, x_2 \dots x_T) \in \Delta_m^T} \frac{1}{T} \sum_{t=1}^T u_O(x_t, y_t)$$

Where  $y_t = \mathcal{A}(x_1, x_2 \dots x_{t-1})$

# Model

## Learner Payoff

With full information (payoffs, learner algorithm), the optimizer plays a best-response sequence of actions


$$x_1, x_2 \dots x_T \in \operatorname{argmax}_{(x_1, x_2 \dots x_T) \in \Delta_m^T} \frac{1}{T} \sum_{t=1}^T u_O(x_t, y_t)$$

Where  $y_t = \mathcal{A}(x_1, x_2 \dots x_{t-1})$

The learner gets payoff

$$V_L(\mathcal{A}, u_O, T) = \frac{1}{T} \sum_{t=1}^T u_L(x_t, y_t)$$

# Model : The Stackelberg Perspective



The Learner Commits to a Learning Algorithm



The Optimizer plays a best-response sequence

The Learner wants to maximize their resulting payoff

A, Collina, Kearns - Solves the full information version of this problem

Our question - What is a good algorithm for the learning version?


# Pareto-Optimality

## Re-defining optimality over all possible optimizers

A property of algorithms based upon a partial order over algorithms. Two Algorithms A and B are compared over all possible optimizer payoffs

A property of algorithms based upon a partial order over algorithms. Two Algorithms \$A\$ and \$B\$ are compared over all possible optimizer payoffs

Three Scenarios:



█ The algorithms do equally well

█ Algorithm A does better

█ Algorithm B does better

# Pareto-Optimality

Re-defining optimality over all possible optimizers

Algorithm A Pareto-dominates algorithm B for some payoff  $u_L$  if:

1.  $\forall \mu_O : V_L(A, u_O) \geq V_L(B, u_O)$
2.  $\exists \mu_O \text{ s.t. } V_L(A, u_O) > V_L(B, u_O)$

An algorithm is Pareto-Optimal if it is not Pareto-dominated

(All results are for positive measure sets and limit average payoffs)

# A Basic Guarantee : No-Regret

Pick action  
 $y_t \in Y$

Get feedback from an adaptive  
adversary.  $f_t : Y \rightarrow [-1,1]$

Objective : Guarantee that the performance is comparable to the single best action in hindsight

$$\text{i.e. } \sum_{t=1}^T f_t(y_t) \geq \max_{y^* \in Y} \sum_{t=1}^T f_t(y^*) - o(T)$$

Realized utility

Best response in hindsight

# No-Regret : Applications

Algorithms exists given convexity

- Online Shortest Path Problem (All s-t paths)
- Online Classification (All classifiers in a concept class)
- Boosting Weak Classifiers (via Minimax Computation)
- Bidders behavior in online auctions is consistent with no-regret learning algorithms [Nekipelov et al., 2015]

# No-Regret : Applications

Algorithms exists given convexity

$$\text{In our setting : } \sum_{t=1}^T u_L(x_t, y_t) \geq \max_{y^* \in \Delta^n} \sum_{t=1}^T u_L(x_t, y^*) - o(T)$$

For example: Rock, Paper and Scissors:

Learner Sequence



Optimizer Sequence



# FTRL

## A Popular class of No-Regret Algorithms

Given that  $R$  is continuous and strongly-convex, and  $\eta_T = \frac{1}{o(T)}$ :

$$y_t = \arg \max_{y \in \Delta^n} \left( \sum_{s=1}^{t-1} u_L(x_s, y) - \frac{R(y)}{\eta_T} \right)$$

All Follow-the-Regularized Leader type algorithms, including Multiplicative Weights (Hedge), Online Gradient Descent  
are Mean-Based No-Regret Algorithms

# A Stronger Guarantee : No-Swap-Regret

Pick action  
 $y_t \in Y$

Get feedback.  $f_t : Y \rightarrow [-1,1]$

Objective : Guarantee that the performance is comparable to any swap function in-hindsight

$$\text{i.e. } \sum_{t=1}^T f_t(y_t) \geq \max_{\pi:Y \rightarrow Y} \sum_{t=1}^T f_t(\pi(y_t)) - o(T)$$

# No-Swap-Regret

A Stronger version of No-Regret

- Calibrated Forecasting
- Boosting for Regression
- Stronger Guarantees exist for context-based subsequences

# Non-Manipulability

The optimizer has an asymptotic best-response that is just playing a static strategy over time

**Is there an algorithm that has all three properties - No-Regret,  
Pareto-Optimality and Non-Manipulability?**

Trivial to achieve any one property

# Our Results

## Main Results

Result 1:

**All No-Swap-Regret algorithms are Pareto-Optimal and non-manipulable.**

PS: The non-manipulability result was already proved by Deng et al. (2019), via a different sequence of arguments

# Our Results

## Main Results

Result 2:

**Not all No-Regret algorithms are Pareto-optimal. Specifically, Follow-the-Regularized-Leader (FTRL) based algorithms (which includes Multiplicative Weights Update, Online Gradient Descent) are Pareto-dominated.**


# Our Results

## Other Results


- A Geometric View of Algorithms
- A characterization of best-responses to mean-based no-regret algorithms (i.e. how to manipulate them)
- A characterization of Pareto-optimal No-Regret Algorithms

# RL Experiment for Optimizer

## Best-Response to Multiplicative Weights




Rock, Paper, Scissors for  $T=1000$




Modified RPS (Non zero sum) for  $T=100$

# RL Experiment for Optimizer

## Best-Response to Multiplicative Weights



Rock, Paper, Scissors for  $T=1000$



Modified RPS (Non zero sum) for  $T=1000$

# Talk Plan

- Geometric View of Learning Algorithms - Menus
- NSR is Non-Manipulable (Intuition)
- FTRL is Pareto-dominated (Intuition) (Time Permitting)
- Future Directions/ Related Work

# Menus

# Summaries of Play

## Transcripts and Correlated Strategy Profiles (CSPs)

### Transcript of Play

Sequence of action pairs

$$\{x_t, y_t\}_{t=1}^T$$

### Correlated Strategy Profile

Empirical distribution over all resulting pure action pairs


$$\phi = \frac{1}{T} \sum_{t=1}^T x_t \otimes y_t$$

CSPs are sufficient to check for no-regret/ no-swap-regret

# CSP : Example

## Transcripts and Correlated Strategy Profiles (CSPs)

Learner Sequence



.....

Algorithm : Mimic the optimizer

Optimizer Sequence



.....

Sequence: Alternate Paper and Rock


$$\text{CSP: } \phi = 1/2(R \otimes P) + 1/2(P \otimes R)$$

# Menus

## All possible CSPs

For every optimizer sequence  $x_1, x_2, \dots, x_T$ , record the induced CSP

Menu of an Algorithm : Take the convex hull of this set



Implicitly : The Limit Menu as  $T \rightarrow \infty$

# Menus: An Example

All possible CSPs

Learning Algorithm A1: Always play P

	A	B
P	X	X
Q	X	X

# Menus: Example 1

All possible CSPs

Learning Algorithm A1: Always play P

	A	B
P	X	X
Q	X	X

$(A \otimes P)$



$(B \otimes P)$



# Menus: An Example

## All possible CSPs

Learning Algorithm A2: Play Q as long as the Optimizer has always played A. Otherwise, play P


	A	B
P	X	X
Q	X	X

# Menus: Example 2

## All possible CSPs


Learning Algorithm A2: Play Q as long as the Optimizer has always played A. Otherwise, play P

	A	B
P	X	X
Q	X	X



# Menus are all you need


## Learner and Optimizer Payoffs



The optimizer “picks” their favorite extreme point

# Menus are all you need

## Pareto-Optimality



# Menus are all you need

## No-Regret : Property of the CSPs

A CSP  $\phi$  is no-regret if, for each  $j \in [n]$ , it satisfies

$$\sum_{i \in [m]} \phi_{ij} u_L(i, j) \geq \max_{j^* \in [n]} \sum_{i \in [m]} \phi_{ij} u_L(i, j^*) .$$

# Menus are all you need

## Non-Manipulability

All Extreme points of the algorithm's menu are product distributions


# Menus are all you need

## Non-Manipulability : A Negative Example

All Extreme points of the algorithm's menu are product distributions

Algorithm : Follow the Leader

Learner Sequence



.....

Optimizer Sequence



.....

Corresponding CSP  $1/3(P \otimes R) + 1/3(S \otimes P) + 1/3(R \otimes S)$

# Menus: Proving Pareto-Optimality

Inclusion-Minimality implies Pareto-Optimality

Every inclusion-minimal menu that contains  $\phi^+$  is Pareto-Optimal.

# Menus: Proving Pareto-Optimality

Inclusion-Minimality implies Pareto-Optimality

Definition: Inclusion-Minimality

A menu  $M_1$  is inclusion-minimal if there is no menu  $M_2$  such that  $M_2 \subsetneq M_1$ .

Definition:  $\varphi^+$

$$\phi^+ = x^* \otimes y^*, \text{ where } (x^*, y^*) = \arg \max_{(x,y)} u_L(x, y)$$

# Menus: Proving Pareto-Optimality

$\phi_+$ -Inclusion-Minimality implies Pareto-Optimality

Lemma : If  $M_1$  contains  $\varphi^+$  and  $M_2 \setminus M_1 \neq \emptyset$ , then there is an Optimizer payoff  $u_O$  such that


$$V_L(M_1, u_O) > V_L(M_2, u_O)$$

Negating Pareto-domination of one algorithm by another requires only a single certificate

# Menus: Proving Pareto-Optimality

Lemma : If  $M_1$  contains  $\varphi^+$  and  $M_2 \setminus M_1 \neq \emptyset$ , then there is an Optimizer payoff  $u_O$  such that


$$V_L(M_1, u_O) > V_L(M_2, u_O)$$



Special case: Both menus are polytopes


# Menus: Proving Pareto-Optimality

Take the convex hull of the union




# Menus: Proving Pareto-Optimality

Start with an “extra” vertex in  $M_2$




# Menus: Proving Pareto-Optimality

1. Start with an “extra” vertex in  $M_2$
2. Construct a path of strictly increasing  $u_L$  value
3. Find a “crossover” edge



# Menus: Proving Pareto-Optimality

1. Start with an “extra” vertex in  $M_2$
2. Construct a path of strictly increasing  $u_L$  value
3. Find a “crossover” edge



# No-Swap-Regret

# No-Swap-Regret

## Characterization (implying Non-Manipulability)

Theorem : The menu of every NSR algorithm is the convex hull of all CSPs of the form  $x \otimes y$ , with  $x \in \Delta_m$  and  
 $y \in \text{BR}_L(x)$


Proof : Via showing that the optimizer always has a static best-response

# No-Swap-Regret

## Characterization (implying Non-Manipulability)

Proof : Via showing that the optimizer always has a static best-response

Consider the optimal transcript, and color based on learner actions



.....




# No-Swap-Regret

## Characterization (implying Non-Manipulability)

Proof : Via showing that the optimizer always has a static best-response

Consider the optimal transcript, and color based on learner actions

Collect all the time steps, by action played (dividing fractionally on steps with mixed strategies)



# No-Swap-Regret

## Characterization (implying Non-Manipulability)


Proof : Via showing that the optimizer always has a static best-response

Collect all the time steps, by action played (dividing fractionally on steps with mixed strategies);

Record the optimizer marginals for each color



$$x_{blue}$$



$$x_{black}$$



$$x_{pink}$$

# No-Swap-Regret

## Characterization (implying Non-Manipulability)


Collect all the time steps, by action played (dividing fractionally on steps with mixed strategies);

Record the optimizer marginals for each color

No-Swap-Regret: Blue, black and pink are respectively best-responses to  $x_{blue}$ ,  $x_{black}$  and  $x_{pink}$



$$x_{blue}$$



$$x_{black}$$



$$x_{pink}$$

# No-Swap-Regret


## Characterization (implying Non-Manipulability)

No-Swap-Regret: Blue, black and pink are respectively best-responses to  $x_{blue}$ ,  $x_{black}$  and  $x_{pink}$

The optimal CSP is now a convex combination of CSPs of the form  $x \otimes y$ , with  $x \in \Delta_m$  and  $y \in BR_L(x)$



$x_{blue}$



$x_{black}$



$x_{pink}$

# No-Swap-Regret

## Characterization (implying Non-Manipulability)

No-Swap-Regret: Blue, black and pink are respectively best-responses to  $x_{blue}$ ,  $x_{black}$  and  $x_{pink}$

The optimal CSP is now a convex combination of CSPs of the form  $x \otimes y$ , with  $x \in \Delta_m$  and  $y \in \text{BR}_L(x)$



Might as well play a single distribution  $x$  and let the NSR learner learn a best-response to  $x$

Reduces to the Stackelberg Equilibrium problem, solvable using  $m$  linear programs

# No-Swap-Regret

## Characterization (proving Pareto-Optimality)

Theorem : The menu of every NSR algorithm is the convex hull of all CSPs of the form  $x \otimes y$ , with  $x \in \Delta_m$  and  $y \in \text{BR}_L(x)$

With a little more effort, we can show that this menu is inclusion-minimal, with some additional characterization of valid menus, i.e., convex sets that can be realized by some learning algorithm.

# FTRL is Pareto-dominated

# Recall : FTRL

Only moves within  $o(T)$  of being the historical best-response action get non-trivial, i.e.,  $\Omega_T(1)$  mass.

Given that  $R$  is continuous and strongly-convex, and  $\eta_T = \frac{1}{o(T)}$ :


$$y_t = \arg \max_{y \in \Delta^n} \left( \sum_{s=1}^{t-1} u_L(x_s, y) - \frac{R(y)}{\eta_T} \right)$$

All Follow-the-Regularized Leader type algorithms, including Multiplicative Weights (Hedge), Online Gradient Descent are Mean-Based No-Regret Algorithms

# Mean-Based Algorithms (FTRL)

Only moves within  $o(T)$  of being the historical best-response action get non-trivial, i.e.,  $\Omega_T(1)$  mass.

Optimizer Sequence



# FTRL is Pareto-dominated

What's the smallest size-game in which we can prove this?

- The optimizer must have more than one action.
- The Learner must have more than 2 actions. Since No-Regret with two actions implies no-swap-regret.

We prove this for a non-degenerate set of  $3 \times 2$  games.


# FTRL is Pareto-dominated

**Theorem:** All FTRL algorithms are Pareto-dominated.

**Proof Sketch:**


1. All FTRL algorithms induce the same menu
2. And the menu is a polytope with a succinct description (implicitly gives the optimizer their exact best response information).

# All FTRL algorithms induce the same menu




# All FTRL algorithms induce the same menu Mean-Based Trajectories

Trajectory has a “clear” leader for all but  $o(T)$  time steps.



# All FTRL algorithms induce the same menu Mean-Based Trajectories

Convert arbitrary trajectories to mean-based trajectories.



# Future Directions

# Future Directions/ Related Work

## Direction 1 : Auctions as repeated Bayesian Games

The learner receives a private context in each round drawn from a prior distribution, for eg., a click through rate prediction for an ad slot

- Mansour et al. (2022) show non-manipulability results via a notion of regret against policies mapping contexts to actions
- Kumar et al. [2024] show similar properties for Online Mirror Descent when used in repeated first price auctions

The Pareto-Optimality question remains open in this setting

# Future Directions/ Related Work

## Direction 2 : Repeated Auctions with a Budget

The learner has a total budget that they can spend, and must optimize spending, possibly based on private contexts

The screenshot shows a website header with 'Meta' and navigation links for 'Research', 'Publications', 'Programs', 'Datasets', and 'Careers'. Below the header, a large title reads 'Pacing Equilibrium in First Price Auction Markets' under the 'Management Science' section.

How do learning  
algorithms for the budget  
pacing problem fare  
against each other?



Thanks for Listening. Questions?